

Adjusted SCM Truncated Regression training

Brian K. Masinde

```
# clearing workspace  
rm(list = ls())
```

```
library(rpart)  
library(dplyr)  
library(caret)  
library(data.table)  
library(mlflow)  
library(reticulate)  
library(Metrics)  
library(themis)  
library(doMC)  
library(here)
```

Inputs

```
# Recipe inputs  
truncated_train <- read.csv(here("data", "truncated_train.csv"))  
truncated_validation <- read.csv(here("data", "truncated_validation.csv"))  
  
df_trunc_train2 <- rbind(truncated_train, truncated_validation)  
  
nrow(df_trunc_train2)
```

```
## [1] 396
```

```
# Fitting tree for wind and rain  
# wind_max prediction using decision trees  
  
trunc_wind_model <- readRDS(here("adjusted SCM/new trunc models", "trunc_wind_model_tuned.rds"))  
  
trunc_rain_model <- readRDS(here("adjusted SCM/new trunc models", "dec_trunc_rain_model_tuned.rds"))
```

Interaction terms (moderators)

```

# Predict using model: "trunc_wind_model"
# To get variable: wind_max_pred

df_trunc_train2[["wind_max_pred"]] <- predict(trunc_wind_model, newdata = df_trunc_train2)

# To get variable: rain_total_pred
df_trunc_train2[["rain_total_pred"]] <- predict(trunc_rain_model, newdata = df_trunc_train2)

# # Define wind and rain interaction variables
wind_fractions <- c("blue_ss_frac", "yellow_ss_frac", "orange_ss_frac", "red_ss_frac")
rain_fractions <- c("blue_ls_frac", "yellow_ls_frac", "orange_ls_frac", "red_ls_frac")

# Compute wind interaction terms dynamically
for (col in wind_fractions) {
  print(col)
  new_col_name <- paste0("wind_", col)
  df_trunc_train2 [[new_col_name]] <- df_trunc_train2 [[col]] * df_trunc_train2 [["wind_max_pred"]]
}

## [1] "blue_ss_frac"
## [1] "yellow_ss_frac"
## [1] "orange_ss_frac"
## [1] "red_ss_frac"

# Multiply rain fractions by rain_total
for (col in rain_fractions) {
  new_col_name <- paste0("rain_", col)
  df_trunc_train2 [[new_col_name]] <- df_trunc_train2 [[col]] * df_trunc_train2 [["rain_total_pred"]]
}

```

Model training

```

# TRUNCATED REGRESSION MODEL TRAINING AND TUNING USING CV
# CV folds and models
n_folds <- 10
n_models <- 25 # adjust depending on search space size, affects seeds length

# Reproducibility: Defining seeds (a little bit complicated because of parallel processing)

# Generate a reproducible list of seeds
set.seed(1234)

seeds_list <- vector(mode = "list", length = n_folds + 1)
for (i in 1:n_folds) {
  seeds_list[[i]] <- sample.int(1000000, n_models) # one seed per model per fold
}
seeds_list[[n_folds + 1]] <- sample.int(1000000, 1) # for final model

```

```

# Set up train control with 10-fold cross-validation
train_control <- trainControl(
  method = "cv",
  number = n_folds,
  summaryFunction = defaultSummary,
  search = "random", # random selection of the expanded grid
  seeds = seeds_list
)

# Detect and register the number of available cores (use all but one)
num_cores <- parallel::detectCores() - 2
registerDoMC(cores = num_cores) # Enable parallel processing

# Measure the time for a code block to run
system.time({
  # Train the model using grid search with 3-fold CV
  trunc_xgb_reg_model <- train(
    damage_perc ~ track_min_dist +
      wind_max_pred +
      rain_total_pred +
      roof_strong_wall_strong +
      roof_strong_wall_light +
      roof_strong_wall_salv +
      roof_light_wall_strong +
      roof_light_wall_light +
      roof_light_wall_salv +
      roof_salv_wall_strong +
      roof_salv_wall_light +
      roof_salv_wall_salv +
      wind_blue_ss_frac +
      wind_yellow_ss_frac +
      wind_orange_ss_frac +
      wind_red_ss_frac +
      rain_blue_ls_frac +
      rain_yellow_ls_frac +
      rain_orange_ls_frac +
      rain_red_ls_frac +
      island_groups, # Confounder adjustment
    data = df_trunc_train2,
    method = "xgbTree",
    trControl = train_control,
    tuneLength = n_models, # this replaces tuneGrid
    metric = "RMSE" # Optimize based on RMSE
  )
  Sys.sleep(2) # This is just an example to simulate a delay
})

```

```

##      user  system elapsed
## 44.659   1.082    7.793

```

```
# Print best parameters
print(trunc_xgb_reg_model$bestTune)
```

```
##   nrounds max_depth      eta   gamma colsample_bytree min_child_weight
## 9      31          7 0.1848925 8.435601         0.6262378          1
##   subsample
## 9 0.9892467
```

Model Logging

```
# Model Logging
```

```
# set tracking URI
mlflow_set_tracking_uri("http://127.0.0.1:5000")
```

```
# Ensure any active run is ended
suppressWarnings(try(mlflow_end_run(), silent = TRUE))
```

```
# Logging metrics for model training and the parameters used
mlflow_set_experiment(experiment_name = "R - SCM - XGBOOST Truncated regression - CV (Training metrics)")
```

```
## [1] "826144865648052556"
```

```
# Ensure that MLflow has only one run. Start MLflow run once.
run_name <- paste("XGBoost Run", Sys.time()) # Unique name using current time
```

```
# Start MLflow run
mlflow_start_run(nested = FALSE)
```

```
## Warning: 'as_integer()' is deprecated as of rlang 0.4.0
## Please use 'vctrs::vec_cast()' instead.
## This warning is displayed once every 8 hours.
```

```
## # A tibble: 1 x 13
##   run_uuid      experiment_id run_name user_id status start_time
##   <chr>          <chr>      <chr>   <chr>   <chr>   <dtm>
## 1 284a2401012a454e8e5~ 826144865648~ blushin~ masinde RUNNI~ 2025-07-24 17:56:28
## # i 7 more variables: artifact_uri <chr>, lifecycle_stage <chr>, run_id <chr>,
## #   end_time <lgl>, metrics <lgl>, params <lgl>, tags <list>
```

```
# Ensure the run ends even if an error occurs
# on.exit(mlflow_end_run(), add = TRUE)
```

```
# ----- best parameters -----
best_params <- trunc_xgb_reg_model$bestTune

# Log each of the best parameters in MLflow
```

```

for (param in names(best_params)) {
  mlflow_log_param(param, best_params[[param]])
}

# ----- train using best parameters
trunc_damage_fit_reg <- train(damage_perc ~ track_min_dist +
                             wind_max_pred +
                             rain_total_pred +
                             roof_strong_wall_strong +
                             roof_strong_wall_light +
                             roof_strong_wall_salv +
                             roof_light_wall_strong +
                             roof_light_wall_light +
                             roof_light_wall_salv +
                             roof_salv_wall_strong +
                             roof_salv_wall_light +
                             roof_salv_wall_salv +
                             wind_blue_ss_frac +
                             wind_yellow_ss_frac +
                             wind_orange_ss_frac +
                             wind_red_ss_frac +
                             rain_blue_ls_frac +
                             rain_yellow_ls_frac +
                             rain_orange_ls_frac +
                             rain_red_ls_frac +
                             island_groups, # Confounder adjustment
                             method = "xgbTree",
                             trControl = trainControl(method = "none"),
                             tuneGrid = best_params, # Use the best parameters here
                             metric = "RMSE",
                             data = df_trunc_train2
                             )

# obtain predicted values
train_predictions <- predict(trunc_damage_fit_reg, newdata = df_trunc_train2)

# Define bin edges
# Define bin edges
bins <- c(0.00009, 1, 10, 50, 100)

# Assign data to bins
bin_labels <- cut(df_trunc_train2$damage_perc, breaks = bins, include.lowest = TRUE, right = TRUE)

# Create a data frame with actual, predicted, and bin labels
data <- data.frame(
  actual = df_trunc_train2$damage_perc,
  predicted = train_predictions,
  bin = bin_labels
)

# Calculate RMSE per bin
unique_bins <- levels(data$bin) # Get unique bin labels

```

```
rmse_by_bin <- data.frame(bin = unique_bins, rmse = NA, count = NA) # Initialize results data frame

for (i in seq_along(unique_bins)) {
  bin_data <- data[data$bin == unique_bins[i], ] # Filter data for the current bin
  rmse_by_bin$rmse[i] <- sqrt(mean((bin_data$actual - bin_data$predicted)^2, na.rm = TRUE)) # Calculate
  rmse_by_bin$count[i] <- nrow(bin_data) # Count observations in the bin
}

# Display RMSE by bin
print(rmse_by_bin)
```

```
##      bin      rmse count
## 1 [9e-05,1]      NaN     0
## 2  (1,10]      NaN     0
## 3  (10,50] 3.750794    322
## 4  (50,100] 7.256758     74
```

```
as.data.frame(rmse_by_bin)
```

```
##      bin      rmse count
## 1 [9e-05,1]      NaN     0
## 2  (1,10]      NaN     0
## 3  (10,50] 3.750794    322
## 4  (50,100] 7.256758     74
```

```
RMSE_1 <- rmse_by_bin[1, "rmse"]
RMSE_10 <- rmse_by_bin[2, "rmse"]
RMSE_50 <- rmse_by_bin[3, "rmse"]
RMSE_100 <- rmse_by_bin[4, "rmse"]

# Log binned RMSE metrics
mlflow_log_metric("RMSE_1", RMSE_1)
```

```
## Warning: 'as_double()' is deprecated as of rlang 0.4.0
## Please use 'vctrs::vec_cast()' instead.
## This warning is displayed once every 8 hours.
```

```
mlflow_log_metric("RMSE_10", RMSE_10)
mlflow_log_metric("RMSE_50", RMSE_50)
mlflow_log_metric("RMSE_100", RMSE_100)
```

```
# End MLflow run
mlflow_end_run()
```

```
## # A tibble: 1 x 13
##   run_uuid      experiment_id run_name user_id status start_time
##   <chr>          <chr>          <chr>   <chr>   <chr>   <dtm>
## 1 284a2401012a454e8e5~ 826144865648~ blushin~ masinde FINIS~ 2025-07-24 17:56:28
## # i 7 more variables: end_time <dtm>, artifact_uri <chr>,
## #   lifecycle_stage <chr>, run_id <chr>, metrics <list>, params <list>,
## #   tags <list>
```

Recipe Outputs

```
# Saving the truncated regression model
full_path <- here("adjusted SCM/new trunc models")

saveRDS(trunc_damage_fit_reg, file = file.path(full_path, paste0("trunc_reg_model", ".rds")))
```

OLD CODE

```
# model_list <- list(
#   wind_max = trunc_wind_model,
#   rain_total = trunc_rain_model,
#   roof_strong_wall_strong = trunc_roof_strong_wall_strong_model,
#   roof_strong_wall_light = trunc_roof_strong_wall_light_model,
#   roof_strong_wall_salv = trunc_roof_strong_wall_salv_model,
#   roof_light_wall_strong = trunc_roof_light_wall_strong_model,
#   roof_light_wall_light = trunc_roof_light_wall_light_model,
#   roof_light_wall_salv = trunc_roof_light_wall_salv_model,
#   roof_salv_wall_strong = trunc_roof_salv_wall_strong_model,
#   roof_salv_wall_light = trunc_roof_salv_wall_light_model,
#   roof_salv_wall_salv = trunc_roof_salv_wall_salv_model
# )
#
# # Apply predictions efficiently
# df_trunc_train2 <- df_trunc_train2 %>%
#   mutate(across(names(model_list), ~ predict(model_list[[cur_column()]],
newdata = df_trunc_train2), .names = "{.col}_pred"))
```