

Model Training: XGBOOST Classifier

Brain K. Masinde

```
# Environment Cleaning: Clearing workspace
rm(list = ls())

# load packages
library(rpart)
library(dplyr)

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(caret)

## Loading required package: ggplot2

## Loading required package: lattice

library(pROC) # For AUC calculation

## Type 'citation("pROC")' for a citation.

##
## Attaching package: 'pROC'

## The following objects are masked from 'package:stats':
##
##   cov, smooth, var

library(data.table)

##
## Attaching package: 'data.table'
```

```
## The following objects are masked from 'package:dplyr':  
##  
##   between, first, last
```

```
library(mlflow)  
library(reticulate)  
library(Matrix)  
library(purrr) # useful for code optimization
```

```
##  
## Attaching package: 'purrr'
```

```
## The following object is masked from 'package:data.table':  
##  
##   transpose
```

```
## The following object is masked from 'package:caret':  
##  
##   lift
```

```
library(themis)
```

```
## Loading required package: recipes
```

```
##  
## Attaching package: 'recipes'
```

```
## The following object is masked from 'package:Matrix':  
##  
##   update
```

```
## The following object is masked from 'package:stats':  
##  
##   step
```

```
library(doMC)
```

```
## Loading required package: foreach
```

```
##  
## Attaching package: 'foreach'
```

```
## The following objects are masked from 'package:purrr':  
##  
##   accumulate, when
```

```
## Loading required package: iterators
```

```
## Loading required package: parallel
```

```
library(here)
```

```
## here() starts at /Users/masinde/Projects/causal_fairness_Ph_IbF
```

Inputs

```
base_train <- read.csv(here("data", "base_train.csv"))
base_validation <- read.csv(here("data", "base_validation.csv"))
base_test <- read.csv(here("data", "base_test.csv"))
```

```
# Combining train and validation datasets to one
# Because we are going to use CV to train the models later
# naming it df_base_train to remain consistent with df naming
df_base_train <- rbind(base_train, base_validation)

cat("number of rows in combined train data:", nrow(df_base_train), sep = " ")
```

```
## number of rows in combined train data: 7184
```

```
# Ensure target variable is a factor
# Ensure the target variable is a factor with valid names
df_base_train$damage_binary_2 <- factor(df_base_train$damage_binary,
                                         levels = c("0", "1"), # Your current levels
                                         labels = c("Damage_below_10", "Damage_above_10")) # New valid l
```

Model Training using CV

```
# Set up train control with custom seeds
n_folds <- 7
n_models <- 25 # adjust depending on search space size, affects seeds length

# Reproducibility: Defining seeds (a little bit complicated because of parallel processing)

# Generate a reproducible list of seeds
set.seed(1234)

seeds_list <- vector(mode = "list", length = n_folds + 1)
for (i in 1:n_folds) {
  seeds_list[[i]] <- sample.int(1000000, n_models) # one seed per model per fold
}
seeds_list[[n_folds + 1]] <- sample.int(1000000, 1) # for final model

# Set up train control with 10-fold cross-validation
```

```

train_control <- trainControl(
  method = "cv",
  number = n_folds, # number of cross-validation folds
  classProbs = TRUE, # Needed for AUC calculation
  summaryFunction = twoClassSummary,
  sampling = "smote", # caret automatically identifies minority class
  search = "random", # random selection of the expanded grid
  seeds = seeds_list, # Ensures reproducibility
)

# Detect and register the number of available cores (use all but one)
num_cores <- parallel::detectCores() - 2
registerDoMC(cores = num_cores) # Enable parallel processing

# Measure the time for a code block to run
system.time({
  # Train the model using grid search with 10-fold CV

  xgb_model <- train(damage_binary_2 ~ track_min_dist +
    wind_max +
    rain_total +
    roof_strong_wall_strong +
    roof_strong_wall_light +
    roof_strong_wall_salv +
    roof_light_wall_strong +
    roof_light_wall_light +
    roof_light_wall_salv +
    roof_salv_wall_strong +
    roof_salv_wall_light +
    roof_salv_wall_salv +
    blue_ss_frac +
    yellow_ss_frac +
    orange_ss_frac +
    red_ss_frac +
    blue_ls_frac +
    yellow_ls_frac +
    orange_ls_frac +
    red_ls_frac,
    data = df_base_train,
    method = "xgbTree",
    trControl = train_control,
    tuneLength = n_models, # this replaces tuneGrid
    metric = "ROC" # "xgbTree" does not support other metrics for classification tasks (e.g., Kappa)
  )
  Sys.sleep(2) # This is just an example to simulate a delay
})

```

```

##      user  system elapsed
## 517.038   4.357   73.441

```

```

# Print best parameters
print(xgb_model$bestTune)

```

```
##   nrounds max_depth      eta   gamma colsample_bytree min_child_weight
## 7      785         6 0.2356713 7.027949         0.6720632         10
##   subsample
## 7   0.540222
```

```
# Extract the best parameters (remove AUC column)
```

```
best_params_model <- xgb_model$bestTune
```

```
damage_fit_class_full <- train(damage_binary_2 ~ track_min_dist +
                               wind_max +
                               rain_total +
                               roof_strong_wall_strong +
                               roof_strong_wall_light +
                               roof_strong_wall_salv +
                               roof_light_wall_strong +
                               roof_light_wall_light +
                               roof_light_wall_salv +
                               roof_salv_wall_strong +
                               roof_salv_wall_light +
                               roof_salv_wall_salv +
                               blue_ss_frac +
                               yellow_ss_frac +
                               orange_ss_frac +
                               red_ss_frac +
                               blue_ls_frac +
                               yellow_ls_frac +
                               orange_ls_frac +
                               red_ls_frac,
                               data = df_base_train, # USE TRAINING AND VALIDATION SETS COMBINED
                               method = "xgbTree", # XGBoost method
                               trControl = trainControl(method = "none"), # No automatic validation
                               tuneGrid = best_params_model # USE BEST PARAMETER
                               )
```

```
# Sanity Check
```

```
# testing on the training datasets (training + validation)
```

```
## Outcome prediction on the final_training_df dataset
```

```
## default function predict returns class probabilities (has two columns)
```

```
y_pred <- predict(damage_fit_class_full,
                  newdata = df_base_train)
```

```
levels(y_pred)
```

```
## [1] "Damage_below_10" "Damage_above_10"
```

```
# using table function
```

```
conf_matrix <- confusionMatrix(y_pred,
                               df_base_train$damage_binary_2, # remember to use damage_binary_2
                               positive = "Damage_above_10"
                               )
```

```
conf_matrix
```

```
## Confusion Matrix and Statistics
##
##               Reference
## Prediction      Damage_below_10 Damage_above_10
##   Damage_below_10             6754             138
##   Damage_above_10              36             256
##
##               Accuracy : 0.9758
##               95% CI : (0.972, 0.9792)
##   No Information Rate : 0.9452
##   P-Value [Acc > NIR] : < 2.2e-16
##
##               Kappa : 0.7339
##
##   Mcnemar's Test P-Value : 1.906e-14
##
##               Sensitivity : 0.64975
##               Specificity : 0.99470
##               Pos Pred Value : 0.87671
##               Neg Pred Value : 0.97998
##               Prevalence : 0.05484
##               Detection Rate : 0.03563
##   Detection Prevalence : 0.04065
##               Balanced Accuracy : 0.82222
##
##   'Positive' Class : Damage_above_10
##
```

```
accuracy <- conf_matrix$overall['Accuracy']
cat("test-set accuracy of associational XGBOOST:", accuracy, sep = " ")
```

```
## test-set accuracy of associational XGBOOST: 0.9757795
```

Model - Mlflow

```
# Logging the model and parameter using MLflow

# set tracking URI
mlflow_set_tracking_uri("http://127.0.0.1:5000")

# Ensure any active run is ended
suppressWarnings(try(mlflow_end_run(), silent = TRUE))

# set experiment
# Logging metrics for model training and the parameters used
mlflow_set_experiment(experiment_name = "Attempt 2: R ass-XGBOOST classification - CV (Training metrics)
```

```
## [1] "345587670307584911"
```

```

# Ensure that MLflow has only one run. Start MLflow run once.
run_name <- paste("XGBoost Run", Sys.time()) # Unique name using current time

# Start MLflow run
mlflow_start_run(nested = FALSE)

## Warning: 'as_integer()' is deprecated as of rlang 0.4.0
## Please use 'vctrs::vec_cast()' instead.
## This warning is displayed once every 8 hours.

## # A tibble: 1 x 13
##   run_uuid          experiment_id run_name user_id status start_time
##   <chr>            <chr>          <chr>   <chr>   <chr>   <dtm>
## 1 8e8958d7ecb54f0cbf0~ 345587670307~ shiveri~ masinde RUNNI~ 2025-07-24 14:08:43
## # i 7 more variables: artifact_uri <chr>, lifecycle_stage <chr>, run_id <chr>,
## #   end_time <lgl>, metrics <lgl>, params <lgl>, tags <list>

# Ensure the run ends even if an error occurs
#on.exit(mlflow_end_run(), add = TRUE)

# Extract the best parameters (remove AUC column)
best_params_model <- xgb_model$bestTune

# Log each of the best parameters in MLflow
for (param in names(best_params_model)) {
  mlflow_log_param(param, best_params_model[[param]])
}

# Log the model type as a parameter
mlflow_log_param("model_type", "ass-xgboost-classification")

# summarize results
conf_matrix <- confusionMatrix(y_pred,
                                df_base_train$damage_binary_2,
                                positive = "Damage_above_10"
                                )

# accuracy
accuracy <- conf_matrix$overall['Accuracy']

# Positive class = 1, precision, recall, and F1
# Extract precision, recall, and F1 score
precision <- conf_matrix$byClass['Precision']
recall <- conf_matrix$byClass['Recall']
f1_score <- conf_matrix$byClass['F1']

# Log parameters and metrics
# mlflow_log_param("model_type", "scm-xgboost-classification")
mlflow_log_metric("accuracy", accuracy)

```

```
## Warning: 'as_double()' is deprecated as of rlang 0.4.0
## Please use 'vctrs::vec_cast()' instead.
## This warning is displayed once every 8 hours.
```

```
mlflow_log_metric("F1", f1_score)
mlflow_log_metric("Precision", precision)
mlflow_log_metric("Recall", recall)
```

```
# Save model
saveRDS(model, file = file.path(path_2_folder, "spam_clas_model.rds"))
```

```
# End MLflow run
mlflow_end_run()
```

```
## # A tibble: 1 x 13
##   run_uuid      experiment_id run_name user_id status start_time
##   <chr>          <chr>      <chr>   <chr>   <chr>   <dtm>
## 1 8e8958d7ecb54f0cbf0~ 345587670307~ shiveri~ masinde FINIS~ 2025-07-24 14:08:43
## # i 7 more variables: end_time <dtm>, artifact_uri <chr>,
## #   lifecycle_stage <chr>, run_id <chr>, metrics <list>, params <list>,
## #   tags <list>
```

Model Testing: Testing on out of sample test data.

```
# Testing on test dataset -----
base_test$damage_binary_2 <- factor(base_test$damage_binary,
                                   levels = c("0", "1"), # Your current levels
                                   labels = c("Damage_below_10", "Damage_above_10")) # New valid l
```

```
# predict for damage_binary
# Make probability predictions for classification
y_preds_probs <- predict(damage_fit_class_full,
                        newdata = base_test, type = "prob")[,2] # Probability of class 1
```

```
# AUC
# Compute AUC (better for classification)
auc_value <- auc(roc(base_test$damage_binary_2, y_preds_probs))
```

```
## Setting levels: control = Damage_below_10, case = Damage_above_10
```

```
## Setting direction: controls < cases
```

```
auc_value
```

```
## Area under the curve: 0.936
```



```

## assigning final class based on threshold
threshold = 0.3
y_pred <- ifelse(y_preds_probs > threshold, 1, 0)

y_pred <- factor(y_pred, levels = c("0", "1"), # Your current levels
                 labels = c("Damage_below_10", "Damage_above_10")) # New valid l

# using table function
conf_matrix <- confusionMatrix(as.factor(y_pred),
                               base_test$damage_binary_2,
                               positive = "Damage_above_10"
                               )
print(conf_matrix)

```

```
## Confusion Matrix and Statistics
```

```
##
##               Reference
## Prediction      Damage_below_10 Damage_above_10
##   Damage_below_10             1649             40
##   Damage_above_10              47             61
```

```
##
##               Accuracy : 0.9516
##               95% CI : (0.9406, 0.961)
##   No Information Rate : 0.9438
##   P-Value [Acc > NIR] : 0.08107
```

```
##
##               Kappa : 0.5581
```

```
##   McNemar's Test P-Value : 0.52005
```

```
##
##               Sensitivity : 0.60396
##               Specificity : 0.97229
##   Pos Pred Value : 0.56481
##   Neg Pred Value : 0.97632
##               Prevalence : 0.05620
##   Detection Rate : 0.03395
##   Detection Prevalence : 0.06010
##   Balanced Accuracy : 0.78812
```

```
##
##   'Positive' Class : Damage_above_10
```

```
cat("Precision of Positive class: Damage above 10:", conf_matrix$byClass['Precision'], sep = " ")
```

```
## Precision of Positive class: Damage above 10: 0.5648148
```

```
cat("Recall of postive class: Damage above 10:", conf_matrix$byClass['Recall'], sep = " ")
```

```
## Recall of postive class: Damage above 10: 0.6039604
```

```
cat("F1 score of positive class: Damage above 10:", conf_matrix$byClass['F1'], sep = " ")
```

```
## F1 score of positive class: Damage above 10: 0.5837321
```

Logging Test metrics - Mlflow

```
# set tracking URI
mlflow_set_tracking_uri("http://127.0.0.1:5000")

# Ensure any active run is ended
suppressWarnings(try(mlflow_end_run(), silent = TRUE))

# set experiment
# Logging metrics for model training and the parameters used
mlflow_set_experiment(experiment_name = "Attempt 2: R ass- XGBOOST classification - CV (Test metircs)")

## [1] "244452174007912648"

# Ensure that MLflow has only one run. Start MLflow run once.
run_name <- paste("XGBoost Run", Sys.time()) # Unique name using current time

# Start MLflow run
mlflow_start_run(nested = FALSE)

## # A tibble: 1 x 13
##   run_uuid          experiment_id run_name user_id status start_time
##   <chr>            <chr>          <chr>   <chr>   <chr>   <dtm>
## 1 459c684fb3a44ba4ada~ 244452174007~ gentle~~ masinde RUNNI~ 2025-07-24 14:08:43
## # i 7 more variables: artifact_uri <chr>, lifecycle_stage <chr>, run_id <chr>,
## #   end_time <lgl>, metrics <lgl>, params <lgl>, tags <list>

# Ensure the run ends even if an error occurs
#on.exit(mlflow_end_run(), add = TRUE)

# Extract the best parameters (remove AUC column)
#best_params_model <- best_params %>% # Remove AUC column if present
#   select(-AUC)

parameters_used <- xgb_model$bestTune

# Log each of the best parameters in MLflow
for (param in names(parameters_used)) {
  mlflow_log_param(param, parameters_used[[param]])
}

# Log the model type as a parameter
mlflow_log_param("model_type", "ass-xgboost-classification")

# predicting
# Threshold
threshold = 0.3
y_preds_probs <- predict(damage_fit_class_full, newdata = base_test, type = "prob")[,2] # Probability
y_pred <- ifelse(y_preds_probs > threshold, 1, 0)
```

```

y_pred <- factor(y_pred, levels = c("0", "1"), # Your current levels
                labels = c("Damage_below_10", "Damage_above_10")) # New valid l

# summarize results
conf_matrix <- confusionMatrix(as.factor(y_pred),
                               base_test$damage_binary_2,
                               positive = "Damage_above_10"
                               )

# accuracy
accuracy <- conf_matrix$overall['Accuracy']

# Positive class = 1, precision, recall, and F1
# Extract precision, recall, and F1 score
precision <- conf_matrix$byClass['Precision']
recall <- conf_matrix$byClass['Recall']
f1_score <- conf_matrix$byClass['F1']
auc_value <- auc(roc(base_test$damage_binary_2, y_preds_probs))

## Setting levels: control = Damage_below_10, case = Damage_above_10

## Setting direction: controls < cases

# Log parameters and metrics
# mlflow_log_param("model_type", "scm-xgboost-classification")
mlflow_log_metric("accuracy", accuracy)
mlflow_log_metric("F1", f1_score)
mlflow_log_metric("Precision", precision)
mlflow_log_metric("Recall", recall)
#mlflow_log_metric("AUC", auc_value)

# Save model
#saveRDS(model, file = file.path(path_2_folder, "spam_clas_model.rds"))

# End MLflow run
mlflow_end_run()

## # A tibble: 1 x 13
##   run_uuid      experiment_id run_name user_id status start_time
##   <chr>          <chr>         <chr>   <chr>   <chr>   <dtm>
## 1 459c684fb3a44ba4ada~ 244452174007~ gentle~ masinde FINIS~ 2025-07-24 14:08:43
## # i 7 more variables: end_time <dtm>, artifact_uri <chr>,
## #   lifecycle_stage <chr>, run_id <chr>, metrics <list>, params <list>,
## #   tags <list>

```

Misc. Experiments

```
# test set confusion matrix
conf_matrix
```

```
## Confusion Matrix and Statistics
##
##               Reference
## Prediction      Damage_below_10 Damage_above_10
## Damage_below_10           1649           40
## Damage_above_10            47           61
##
##               Accuracy : 0.9516
##               95% CI : (0.9406, 0.961)
##       No Information Rate : 0.9438
##       P-Value [Acc > NIR] : 0.08107
##
##               Kappa : 0.5581
##
## Mcnemar's Test P-Value : 0.52005
##
##       Sensitivity : 0.60396
##       Specificity : 0.97229
##       Pos Pred Value : 0.56481
##       Neg Pred Value : 0.97632
##       Prevalence : 0.05620
##       Detection Rate : 0.03395
##       Detection Prevalence : 0.06010
##       Balanced Accuracy : 0.78812
##
##       'Positive' Class : Damage_above_10
##
```

```
# get the false positives by regions (Luzon, Visayas, Mindanao)
# confusion matrix by regions
# Make sure the grouping variable is a factor
# Make sure island_groups is a factor
base_test$island_groups <- as.factor(base_test$island_groups)

# Loop through each group and generate a confusion matrix
for (grp in levels(base_test$island_groups)) {

  # Subset data for the current group
  group_indices <- base_test$island_groups == grp
  y_true_group <- base_test$damage_binary_2[group_indices]
  y_pred_group <- y_pred[group_indices]

  # Generate and print confusion matrix
  cat("Confusion Matrix for Island Group:", grp, "\n")
  print(confusionMatrix(y_pred_group, y_true_group, positive = "Damage_above_10"))
  cat("\n")
}
```

```
## Confusion Matrix for Island Group: Luzon
```

```

## Confusion Matrix and Statistics
##
##               Reference
## Prediction      Damage_below_10  Damage_above_10
##   Damage_below_10              1188             28
##   Damage_above_10               22             30
##
##               Accuracy : 0.9606
##               95% CI : (0.9483, 0.9706)
##   No Information Rate : 0.9543
##   P-Value [Acc > NIR] : 0.1565
##
##               Kappa : 0.5249
##
##   McNemar's Test P-Value : 0.4795
##
##               Sensitivity : 0.51724
##               Specificity : 0.98182
##   Pos Pred Value : 0.57692
##   Neg Pred Value : 0.97697
##   Prevalence : 0.04574
##   Detection Rate : 0.02366
##   Detection Prevalence : 0.04101
##   Balanced Accuracy : 0.74953
##
##   'Positive' Class : Damage_above_10
##
## Confusion Matrix for Island Group: Mindanao
## Confusion Matrix and Statistics
##
##               Reference
## Prediction      Damage_below_10  Damage_above_10
##   Damage_below_10              103             3
##   Damage_above_10               1             3
##
##               Accuracy : 0.9636
##               95% CI : (0.9095, 0.99)
##   No Information Rate : 0.9455
##   P-Value [Acc > NIR] : 0.2775
##
##               Kappa : 0.5817
##
##   McNemar's Test P-Value : 0.6171
##
##               Sensitivity : 0.50000
##               Specificity : 0.99038
##   Pos Pred Value : 0.75000
##   Neg Pred Value : 0.97170
##   Prevalence : 0.05455
##   Detection Rate : 0.02727
##   Detection Prevalence : 0.03636
##   Balanced Accuracy : 0.74519
##

```

```

##          'Positive' Class : Damage_above_10
##
##
## Confusion Matrix for Island Group: Visayas
## Confusion Matrix and Statistics
##
##              Reference
## Prediction      Damage_below_10  Damage_above_10
##   Damage_below_10              358              9
##   Damage_above_10              24              28
##
##              Accuracy : 0.9212
##              95% CI : (0.8912, 0.9452)
##   No Information Rate : 0.9117
##   P-Value [Acc > NIR] : 0.27848
##
##              Kappa : 0.5866
##
## Mcnemar's Test P-Value : 0.01481
##
##              Sensitivity : 0.75676
##              Specificity : 0.93717
##              Pos Pred Value : 0.53846
##              Neg Pred Value : 0.97548
##              Prevalence : 0.08831
##              Detection Rate : 0.06683
##   Detection Prevalence : 0.12411
##   Balanced Accuracy : 0.84696
##
##          'Positive' Class : Damage_above_10
##

```

Outputs

```

saveRDS(damage_fit_class_full, here("associational XGBOOST", "ass_XGBOOST_class.rds"))

```

OLD CODE